

Piazza



Enabling Government Teams to Share and Access Data in the Cloud in 2016

Michael P. Gerlek
mgerlek@radiantblue.com

4 May 2016

*"A **piazza** is commonly found at the
meeting of two or more streets...*

*Shops and other small businesses
are found on piazzas...*

*Metro stations and bus stops
are found on piazzas...*

*An ideal place to set up
a business..."*

*"An infrastructure that
enables geospatial services"*

Hello, Piazza!

in which we attempt to justify our existence

Within the government, we see:

- Teams of analysts trying to solve problems
- Often limited by
 - Access to new data sources
 - Access to new tools, technologies
 - Access to scalable hardware, ...
- And often only a handful of supporting devs
 - JavaScript devs, Python data scientists, ...

➔ Our target customers

And they're all under orders to

Move To The Cloud!

Make the "easy" things easy
And let the hard thing be hard

We can do the heavy lifting
Because you shouldn't have to

*(Because many of you
don't need all the hard bits)*

- Data access
 - *Where is it? How can I access it? Which ones have kittens?*
- Workflows
 - *Do X, and then when Y happens, go do Z.*
- User services
 - *Go run my new, cool algorithm! At scale!*
- Security, Auditing, Logging
 - *Hands off my stuff.*

- Build a set of web services
 - All cloud-native
 - All open source
- Providing essential functionality
 - Load, access, search
 - Orchestrate, execute
- All under one framework
 - Without exposing any of the complexity

An Example

in which we attempt, again, to justify our
existence

I HAVE A FRIEND...

- **The Shoreline Extraction Problem**
 - Given a set of (coastal) images
 - Compute shoreline vectors
- *Every N years, do:*
 - Collect a big pile of imagery data
 - Hire a bunch of contractors
 - Draw coastlines, insert into database
 - (Discard all intermediate products, supporting scripts, and process documentation)

- As the night follows the day...
 - Better shoreline detection algorithms are written
 - More and more imagery is collected
 - People start asking for up-to-date shoreline data in their AOIs
- So **someone** has to automate this process
 - (And that **someone** is not a rock star hipster geospatial dev)

And so *someone* needs to:

- Harvest metadata from large datasets
- Search for AOIs in all that data
- React when new imagery becomes available
- Run the detection algorithms

Oh, and:

- Do everything in the cloud
- Do everything at scale
- Do everything automatically

The Features

in which we show what Piazza can do

- The “No-Host” model
 - Piazza is not intended to be a data hub
 - Rather, Piazza is a proxy of sorts
- Registration and Metadata
 - URL of data source
 - Title, bounding box, fitness-for-use, ...
 - Features, images, point clouds
- Resource IDs

N x M?
N + M!



https://raw.githubusercontent.com/venicegeo/pz-gateway/

api_key

Explore

Gateway API

Piazza Core Services API

Created by Patrick Doody

See more at <https://github.com/venicegeo/venice/wiki>

[Contact the developer](#)

Alert

Show/Hide

List Operations

Expand Operations

Data

Show/Hide

List Operations

Expand Operations

GET	/data	Query Piazza Data
POST	/data	Load Data into Piazza
POST	/data/file	Load a Data File into Piazza
GET	/data/{dataId}	Get Metadata for Loaded Data
POST	/data/query	Query Metadata in Piazza Data holdings
POST	/deployment	Obtain a GeoServer deployment for a Data Resource object.
GET	/file/{dataId}	Download Data File

Deployment

Show/Hide

List Operations

Expand Operations

Event

Show/Hide

List Operations

Expand Operations


```
{  type: ingest,
  host: true,
  data: {
    dataType: { type: raster },
    metadata: {
      name: test.tif,
      description: foss4g_test
    }
  }
...}
```

```
curl -X POST ...
test.tif ...
https://pz-gateway/data/file
```

→ *cd504e20-cb90-4ff3-bd4c-f755239f2bfd*

- Get Metadata
 - System-extracted, user-supplied
- Request download link
 - “export as”
- Request WMS, WFS layers
 - On the fly, via GeoServer, with leasing

```
curl -X GET ...  
https://pz-gateway/data/cd504e...
```

- Indexing
 - Metadata extracted during data load
 - File format parsing
 - User-supplied fields
- Querying
 - Elasticsearch DSL (for now)

POST /data/query

Users want to call their own algorithms

- With job management, with scaling

- Registration

- URL to web API

- Description of parameters

- Metadata

- **POST /service**

- service ID

- Invocation/execution
 - Service ID
 - Parameters
 - **POST /job**
 - job ID
- Status
 - Progress, results
 - **GET /job**
 - resource ID of result data

Remember that shoreline example?

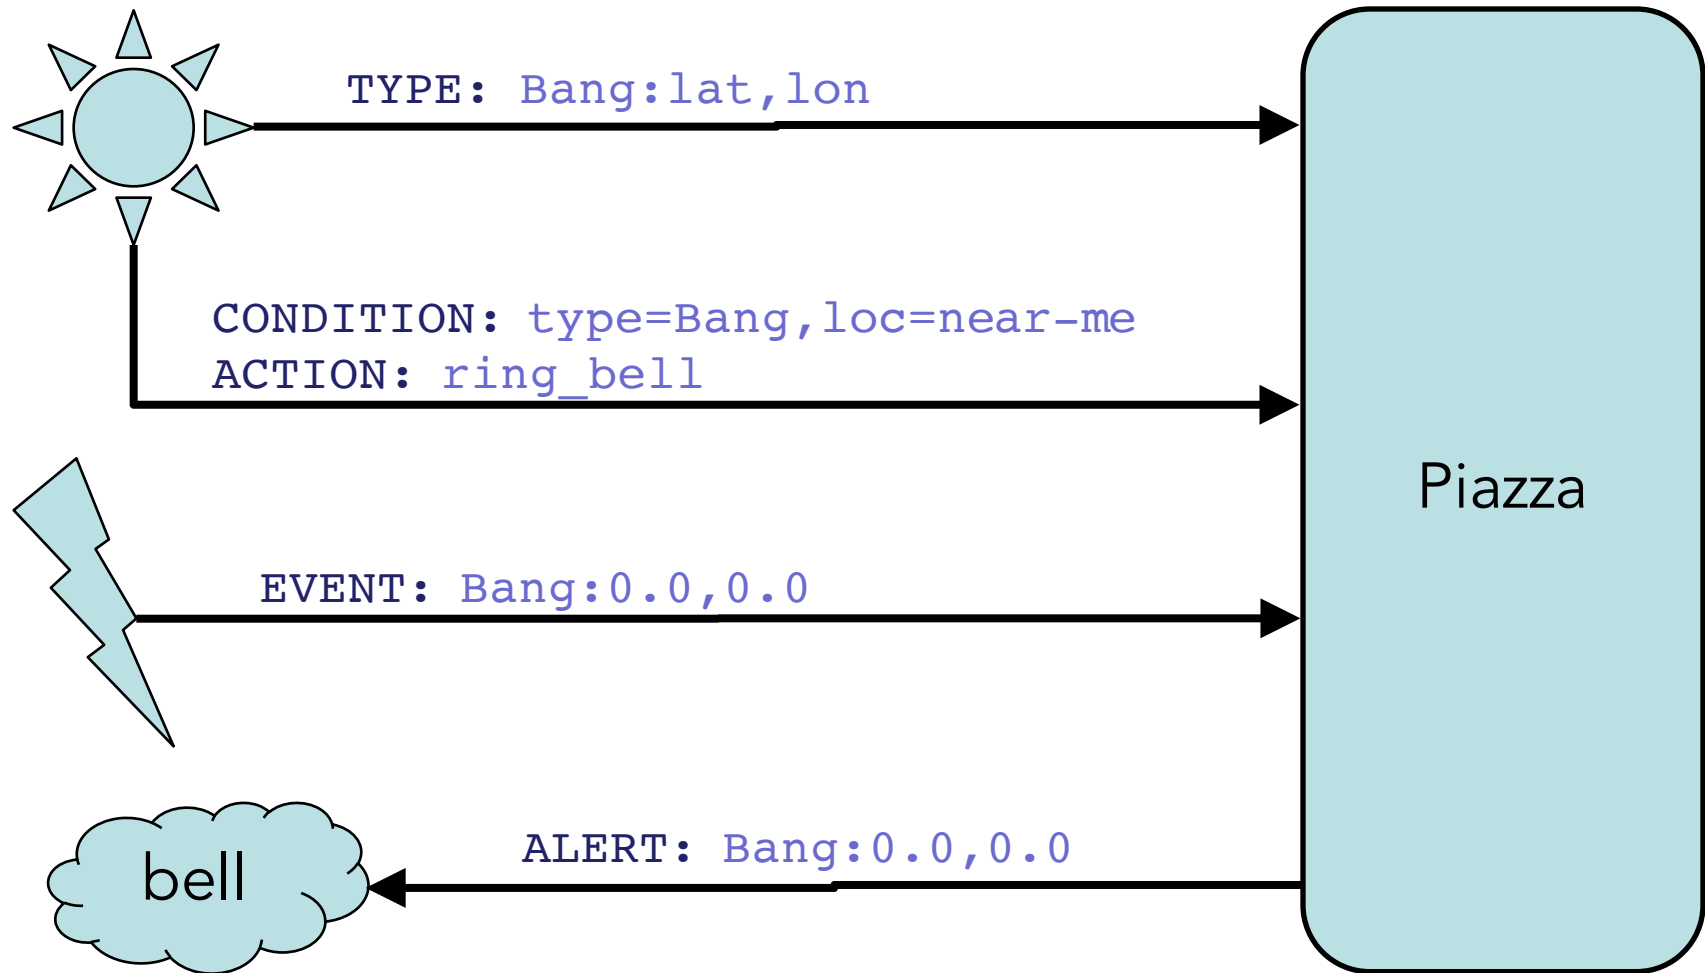
- Watch for new data
- Run an algorithm
- On analysis results:
 - If good, insert to database
 - If bad, push to queue for manual checks

"Something happened!"

- Name, ID
- List of parameters
- *System-level*
 - e.g. "image was loaded"
 - Issued by Piazza's internal services
- *User-level*
 - e.g. "interesting new data from my sensor"
 - User-defined parameter list
 - POST /eventType
 - Issued by some external entity, i.e. client-side
 - POST /event

- The Condition
 - If *eventType* = ...
 - And parameter \geq ...
- The Action
 - Invoke some service
 - With parameter substitution!
 - Simplest: post to Alert Queue

Think IFTTT



Architecture

in which we provide the obligatory
boxes-and-lines diagrams

Have you had to work with a complex library controlled by an equally complicated API?

Have you had to work with a complex library controlled by an equally complicated API?


And found yourself writing a simpler API layer just to hide all that complexity?

Have you had to work with a complex library controlled by an equally complicated API?

And found yourself writing a simpler API layer just to hide all that complexity?

Hi. Me too.

THE BLACK BOX



WHAT PIAZZA IS *Not*

- An **application**
 - It's a platform/architecture/framework/whatever
 - (Indeed, it's pretty useless by itself)
- A **universal solution** to everyone's problems
 - Will always need custom systems
- A "**replacement**" for anything
 - Focus is on new needs, new workflows
- Any **new technology** or **rocket science**
 - Use existing technologies and best practices

Client App

...

Piazza

...

Cloud Foundry

...

AWS

- We build on top of Cloud Foundry
 - *For now, anyway*
- Standing up Piazza is nontrivial
 - GeoServer, PostGIS, Elasticsearch, ...
 - 20+ microservices
- Which is okay(ish) if Piazza is hosted for you
 - But will make you very sad otherwise

We have a lot of work to do here.

The Next Steps

in which we preview next year's FOSS4G talk

- User services
 - URLs should really be deployable objects
 - It's Piazza's job to stand up and scale
- Load-time services
 - User-supplied file formats, feature filters, ...
- Platform deployability
 - OpenShift, et al
- Enterprise-level cataloging
 - Harvesting metadata in bulk
 - Standards #makeitstop #justshootmenow

- Security
 - *Umm, yeah. That.*
- Developer docs
- *And of course*
 - *Better use cases*
 - *Better users*
 - *A community*

<a href/>

pz-docs.venicegeo.io
/userguide

pz-swagger.venicegeo.io




Prominent Edge



(AND MANY MORE)

Thank you.

Questions?



Piazza